

Fleet-Wide Condition Monitoring via Shared Nominal Representations and Local Radius Calibration

Enabling OEM-Centralized, Fleet-Deployed CBM

Eduardo Di Santi¹

¹Department of Applied Mathematics, University of Colorado Boulder

Abstract

Objectives: Deploying condition monitoring across large fleets of heterogeneous assets remains challenging because representation learning, threshold tuning, and fault-labelled commissioning are typically performed independently for every monitored unit. This work investigates whether these tasks can be decoupled by sharing a single nominal representation across the fleet while allowing each asset to identify only its own local operational radius during commissioning.

Methods: A canonical nominal representation, implemented as a Conv1D autoencoder, is trained once offline using simulated healthy signals and transferred unchanged to every asset. Each deployed unit estimates its operational radius as an empirical residual quantile from a short stream of unlabelled nominal observations using a convergence-based commissioning procedure. The methodology is evaluated under controlled operating conditions using the Case Western Reserve University (CWRU) bearing dataset and subsequently validated on four independent run-to-failure bearings from the NASA IMS dataset. Results are compared with complete-data oracle calibration, direct simulator-threshold transfer, and autoencoders trained independently on the early-life data of each NASA bearing.

Findings: Direct transfer of the absolute simulator threshold produced complete failure after simulation-to-field transfer, confirming that the transferable object is the nominal representation rather than its absolute reconstruction-error scale. Across all CWRU operating conditions, convergence-based commissioning achieved 100% calibration success, an average post-calibration healthy false-alarm rate of 3.59%, ball-fault detection of 99.69%, and 100% detection of inner- and outer-race faults, while recovering operational radii within 2.08% of the complete-data oracle. On NASA IMS, independently commissioned assets converged to asset-specific operational radii after approximately 50 recordings, identified persistent degradation in all four bearings, and provided an average warning of 45.4 hours before the documented end of the experiment, without requiring per-asset representation learning.

Novelty: The contribution is a fleet deployment methodology rather than a new anomaly-detection architecture. The deployable model is the shared nominal representation; deployment of a new asset requires only estimating its local operational radius during commissioning rather than training a new representation or engineering an asset-specific threshold.

Practical Applications: A single validated nominal representation can be developed, versioned, and maintained centrally, for example by an original equipment manufacturer, and deployed unchanged across an operator’s fleet. Each installed asset requires only a short unsupervised commissioning phase to estimate its local operational radius, eliminating per-asset representation training and asset-specific threshold engineering.

The deployable model is the shared nominal representation; deployment of a new asset requires only estimating its local operational radius rather than training a new model.

Keywords: condition-based maintenance (CBM), rail asset management, fleet-wide deployment, asset commissioning, bearing condition monitoring, OEM deployment, shared nominal representation, simulation-to-field transfer, edge AI, run-to-failure validation

1 Introduction

Modern transportation systems increasingly rely on condition-monitoring technologies to supervise large populations of operational assets, including rolling-element bearings, axleboxes, traction motors, gearboxes, pumps, actuators, vehicles, and spatially indexed infrastructure segments Jardine et al. 2006; Zhao et al. 2019. Rolling-element bearings are among the most safety-critical of these components: railway axlebox bearings alone are monitored across thousands of independently operating vehicles, where defects can propagate to derailment-level failures if undetected Entezami et al. 2020. Although machine-learning research has produced increasingly capable anomaly-detection models Chandola et al. 2009; Zhao et al. 2019, the deployment problem remains largely unresolved: a detector that performs well for one asset under one calibrated operating condition does not automatically scale to hundreds or thousands of independently commissioned units.

In practice, fleet growth therefore scales not only the number of monitored assets but also the engineering effort required to train, validate, version, and maintain asset-specific models.

At fleet rollout, several constraints arise simultaneously. First, representative fault examples are rarely available. Second, waiting months or years for failures to accumulate is operationally unacceptable. Third, installation-specific factors such as sensor coupling, operating point, mounting, and environmental noise alter the numerical scale of anomaly scores between assets Randall and Antoni 2011. In railway fleets, for example, axlebox bearings of nominally identical rolling stock are exposed to different track quality, wheel wear, load distribution, and wayside versus on-board sensor placement, so that a detector calibrated on one vehicle rarely transfers its absolute decision threshold to another Entezami et al. 2020. Finally, training, calibrating, validating, and maintaining a separate representation model for every physical unit creates an engineering burden that grows directly with fleet size.

The practical question is therefore not simply

Q1. *How can an anomalous condition be detected?*

but rather

Q2. *How can a single nominal representation be deployed across a heterogeneous fleet while allowing every physical asset to identify only its own local operational neighbourhood, without per-asset representation learning or manual threshold engineering?*

This paper addresses Q2 by separating representation engineering from deployment calibration. A canonical nominal representation is learned once, offline, from simulated healthy observations and subsequently deployed unchanged across the fleet. During commissioning, each physical asset identifies only its own operational radius from a short stream of unlabeled nominal observations. The shared representation remains fixed throughout the asset lifetime.

This separation also defines an OEM-centralized deployment model. The shared nominal representation is engineered, validated, versioned, and maintained centrally, for example by the original equipment manufacturer (OEM), whereas each deployed asset performs only a lightweight local commissioning. Consequently, fleet expansion no longer requires developing, validating, and maintaining an independent representation model for every physical unit. Instead, deployment consists only of distributing the shared representation and estimating one local operational radius for each asset.

From a broader perspective, deployment can be interpreted as a local inverse problem. The shared representation is recovered once during offline model engineering, whereas commissioning consists only of identifying the local operational neighbourhood of each physical asset. In this sense, deployment does not learn a new representation; it estimates one physical property of the monitored asset. This interpretation follows the inverse-problem and quotient-space formulations introduced in our previous work Di Santi 2026a,b, while the present paper addresses their engineering realization for scalable fleet deployment.

The transferable component is therefore not an anomaly threshold but the shared nominal representation itself. Absolute reconstruction-error thresholds generally do not transfer because their numerical scale depends on the simulator, preprocessing, operating conditions, and residual statistics. Instead, the proposed methodology transfers one shared nominal representation together with a fleet-wide nominal coverage, while each asset independently estimates the local operational radius required to realise that coverage. The deployable model is therefore common to the fleet; only the operational radius is asset-specific.




Figure 1: Conventional fleet deployment requires engineering, validating, versioning and maintaining an independent nominal representation for each monitored asset. The proposed deployment architecture instead distributes one shared nominal representation unchanged across the fleet, while each asset estimates only its local operational radius during commissioning. The NASA IMS experiments demonstrate that these locally estimated radii differ substantially between physical assets, despite using the same representation, the same target nominal coverage, and the same convergence criterion.

This formulation also changes the relevant experimental comparison. The appropriate reference is not another globally shared detector but an operationally unrealistic upper bound in which every physical asset has its complete nominal history available for independent representation learning and calibration. Evaluating Q2 therefore requires answering a third question:

Q3. *What is the operational cost of replacing independently engineered asset-specific representations with one centrally maintained shared representation and lightweight local commissioning?*

Q3 is evaluated in two complementary deployment scenarios. The first uses the CWRU Smith and Randall 2015 benchmark to quantify simulation-to-field transfer under controlled operating-condition variability while comparing the proposed methodology with ideal asset-specific representations. The second uses the NASA IMS run-to-failure dataset to evaluate the complete deployment protocol on independently monitored physical assets undergoing naturally evolving degradation.

The asset-specific reference methods are trained and calibrated using the complete healthy dataset of each operating condition. In contrast, the proposed method uses no field observations for representation learning and estimates the local operational radius from only the first 100 healthy windows, corresponding to approximately ten seconds of signal.

The experimental results show that the cost of scalable deployment is small. Under controlled CWRU operating conditions, replacing asset-specific representations with one shared simulator-trained representation reduced ball-fault detection only marginally (99.69% versus 100%) while preserving complete inner- and outer-race detection. The local operational radius converged after approximately fifty commissioning observations without requiring any field retraining.

The NASA IMS experiments further demonstrate that the same deployment principle extends to independently monitored physical assets. All four bearings identified distinct local operational

radii using the same shared representation, the same target nominal coverage, and the same convergence criterion. After commissioning, every bearing subsequently departed persistently from its own estimated nominal neighbourhood, providing early warning between 14.5 and 74 hours before the documented end of the experiment and supporting the interpretation of the operational radius as a physical property of the local nominal behaviour rather than a tuned decision threshold.

The contribution of this paper is therefore not a new anomaly-detection architecture but a deployment methodology for fleet-scale condition monitoring. Although a Conv1D autoencoder is used as the experimental representation, the deployment principle is representation-agnostic. The principal contributions are:

1. an OEM-centralized fleet-deployment architecture in which a single shared canonical nominal representation is engineered, validated, and maintained centrally, while deployed assets perform only lightweight local operational-radius calibration;
2. a direct simulation-to-field deployment methodology in which a single shared nominal representation is transferred unchanged across heterogeneous assets without field-data retraining;
3. an online commissioning procedure that estimates each asset’s local operational radius from a short stream of unlabeled observations;
4. a convergence criterion that determines automatically when sufficient nominal behaviour has been observed to complete commissioning and freeze the operational radius;
5. experimental evidence that independently monitored physical assets identify substantially different local operational radii while sharing the same representation, target nominal coverage, and convergence criterion;
6. an evaluation against ideal asset-specific detectors trained and calibrated with complete healthy datasets; and
7. an edge-oriented deployment model that eliminates per-asset representation learning after installation, allowing fleet-wide deployment through distribution of one shared representation and local unsupervised calibration.

The remainder of the paper is organized as follows. Section 2 reviews condition monitoring, one-class anomaly detection, adaptive calibration, and fleet deployment. Section 3 introduces the deployment methodology. Section 4 presents the experimental protocol. Section 5 reports the deployment results. Section 6 summarizes the mathematical interpretation underlying the deployment guarantees. Section 7 closes with a discussion of industrial implications, limitations, and future work.

2 Related Work

Condition-based maintenance (CBM) has become a central component of modern transportation asset management, where condition monitoring supports the operation of large fleets of heterogeneous assets Chandola et al. 2009; Jardine et al. 2006; Zhao et al. 2019. Considerable progress has been made in improving anomaly-detection algorithms; however, comparatively little attention has been paid to how these detectors are engineered, commissioned, and maintained at fleet scale once thousands of independent physical assets must be deployed.

Most machine-learning approaches formulate this problem as supervised fault classification, requiring representative fault examples during training. While effective on benchmark datasets, these methods assume that deployment conditions remain close to those observed during model development.

One-class methods, including one-class SVM Schölkopf et al. 1999, neural-network novelty detection Bishop 1994, and autoencoder-based anomaly detection Hinton and Salakhutdinov 2006; Sakurada and Yairi 2014, reduce the dependence on fault labels by learning only nominal behaviour. More recent representation-learning approaches further improve anomaly detection by constructing latent spaces in which healthy operation occupies a compact region while degraded behaviour becomes more easily separable.

Most existing work therefore focuses on representation learning itself. Comparatively little attention has been devoted to the deployment architecture: how a validated nominal representation can be commissioned, governed, and reused across an entire fleet while allowing each physical asset to retain its own operational characteristics.

Adaptive thresholds have been widely used to compensate for changing operating conditions. However, unrestricted adaptation risks absorbing progressive degradation into the nominal baseline. This motivates separating calibration from monitoring. In the proposed methodology, adaptation is limited to a finite cold-start phase during which each asset estimates its local operational radius. Once convergence is achieved, the shared representation remains fixed and subsequent departures become diagnostic evidence rather than input for further adaptation.

The proposed deployment strategy is also related to invariant representation learning Bronstein et al. 2021 and anomaly detection under distribution shift Carvalho et al. 2023, both of which seek representations that are insensitive to nuisance variability. Our contribution, however, lies at a different level. Rather than proposing a new representation-learning architecture, we address the deployment problem by combining a shared nominal representation with local operational-radius calibration, eliminating per-asset representation learning while preserving a common nominal coverage across the fleet.

A related and rapidly growing literature addresses cross-condition and cross-machine transfer through domain adaptation, typically aligning source and target feature distributions via maximum mean discrepancy, adversarial domain classifiers, or pseudo-labeled fine-tuning Chen et al. 2026; Hue et al. 2025; Wang et al. 2026. These methods substantially reduce the need for labeled target-domain faults and have demonstrated strong cross-condition accuracy on CWRU and related benchmarks. However, they still require access to target-domain observations during an explicit adaptation or fine-tuning stage, executed independently for each new operating condition or physical unit. This reintroduces, at a smaller scale, the same per-asset engineering burden that fleet-wide deployment seeks to eliminate: every newly installed asset still triggers a training procedure, even when it requires no fault labels.

The distinction is therefore not between supervised and unsupervised deployment, nor between labeled and unlabeled adaptation, but between architectures that require engineering work whenever a new asset is deployed and one in which deployment consists only of estimating a local operational radius while reusing an already validated shared representation.

From an engineering perspective, this also changes the transferable object. Existing transfer-learning approaches transfer model parameters and then adapt them to each target asset. In the proposed methodology, the transferable engineering artefact is the validated nominal representation itself, whereas the only asset-specific quantity estimated during deployment is the local operational radius.

This viewpoint is also consistent with the inverse-problem and quotient-space formulations introduced in our previous work Di Santi 2026a,b, in which diagnosis is posed as the recovery of a canonical nominal representation from observations. The present paper addresses the engineering realization of that framework for scalable fleet deployment rather than its mathematical formulation.

3 Method

This section formalizes the proposed deployment methodology. The central idea is to separate representation engineering from deployment commissioning. A single nominal representation is engineered once and shared across the fleet, whereas each deployed asset identifies only its own local operational radius during commissioning.

3.1 Deployment architecture: shared representation, local operational radius

Let $x_{k,t} \in \mathcal{X}$ denote an observation acquired from operational asset k at time t . In the present experiments, $x_{k,t}$ corresponds to a fixed-length vibration recording.

A shared nominal representation, trained once offline, maps each observation to a scalar residual score

$$r_{k,t} = r(x_{k,t}). \quad (1)$$

During commissioning, each deployed asset estimates its own local operational radius, denoted by τ_k^* , from a short sequence of unlabeled nominal observations. This radius characterizes the local neighbourhood of the nominal manifold represented by the shared model. Once commissioning converges, the operational radius is frozen and remains fixed throughout subsequent monitoring.

An alarm is therefore generated whenever

$$r_{k,t} > \tau_k^*. \quad (2)$$

The deployment methodology follows three principles:

1. one centrally engineered, validated, versioned, and governed nominal representation;
2. one lightweight commissioning procedure performed independently by each deployed asset;
3. no representation retraining, fine-tuning, or parameter adaptation after deployment.

This separation directly matches the OEM-centred deployment model introduced in Section 1. The shared representation is engineered, validated, versioned, and maintained centrally by whichever organization owns the diagnostic model, for example the component OEM, a third-party analytics provider, or the fleet operator itself. Each deployed asset performs only a lightweight local commissioning procedure in which its own operational radius is identified automatically from nominal observations.

Consequently, fleet growth no longer requires engineering additional representation models. Deploying a new physical asset consists only of distributing the already validated shared representation and estimating one local operational radius for that asset.

The transferable engineering artefact is therefore the nominal representation itself rather than an absolute decision threshold. Different assets identify different operational radii while sharing exactly the same representation, the same target nominal coverage, and the same commissioning procedure.

4 Experimental Design

The experimental protocol evaluates the proposed deployment methodology rather than the underlying representation-learning architecture. Across the five experimental stages, we measure:

- false-alarm rate during cold start;
- number of observations required for operational-radius convergence;

- variability of τ_k^* across assets and operating regimes;
- post-convergence fault detection performance;
- sensitivity to convergence parameters;
- agreement between online and complete-data calibration; and
- deployment cost avoided by eliminating per-asset representation learning.

Stage 1 — Controlled multi-asset simulation. Four bearing streams are generated with a physics-motivated nominal simulator using 1200-sample windows, a 30 Hz shaft frequency, and bounded variation in amplitude, phase, speed, offset, and additive noise. The shared representation is trained once on 4,000 simulated healthy windows, with 1,000 additional nominal windows reserved for validation. Faults are introduced only during evaluation. A separate synthetic test set contains 750 windows per operating regime.

Each simulated asset begins in nominal operation and independently estimates its operational radius. Asset A remains nominal, whereas Assets B–D transition after convergence to outer-race, inner-race, and ball-fault regimes.

This stage measures

- (i) observations required for convergence;
- (ii) false alarms before and after convergence;
- (iii) separation between reconstruction residuals and the converged operational radius after fault onset; and
- (iv) whether independently deployed assets converge to stable but asset-specific operational radii.

Stage 2 — Direct simulation-to-field deployment on CWRU. The nominal representation trained in Stage 1 is transferred unchanged to the Case Western Reserve University (CWRU) Drive-End dataset. No CWRU observation is used for representation learning, retraining, or fine tuning.

Each operating condition is treated as an independent deployment environment. During commissioning, healthy observations are processed sequentially. Calibration begins after a minimum of $N_{\min}$ healthy windows and continues until either

- (i) the operational radius satisfies the convergence criterion, or
- (ii) a maximum calibration budget $N_{\max}$ is reached.

Once convergence is declared, the operational radius is frozen and all subsequent observations are evaluated using the unchanged shared representation.

This stage evaluates whether

- (i) the learned representation transfers directly from simulation to field;
- (ii) local empirical calibration compensates for the non-transferability of absolute reconstruction-error thresholds; and
- (iii) the calibration duration can be determined automatically from the observed convergence behaviour rather than prescribed a priori.

The four CWRU operating conditions provide a controlled evaluation of score-scale heterogeneity and deployment calibration rather than four physically distinct fleet assets. Validation on independently monitored run-to-failure bearings is addressed in Stage 4 using the NASA IMS dataset.

Stage 3 — Reference comparison and calibration ablation. To quantify the cost of scalability, we construct an intentionally favorable but operationally non-scalable reference condition. For each CWRU operating condition, three detectors are trained and calibrated independently using the complete healthy dataset:

- (i) a Conv1D autoencoder;
- (ii) a one-class SVM operating on PCA-reduced standardized windows; and
- (iii) an Elliptic Envelope operating on PCA-reduced standardized windows.

Each detector uses the empirical 98th percentile of its own healthy anomaly-score distribution as its decision boundary. Fault observations are never used for model fitting or threshold selection.

Using the unchanged simulator-trained representation, we compare three deployment strategies:

1. **Direct threshold transfer:** apply the absolute simulator threshold unchanged to each CWRU operating condition;
2. **Field oracle:** estimate the operational radius from the complete healthy dataset available for each operating condition, without modifying the shared representation;
3. **Proposed online calibration:** estimate the operational radius sequentially from healthy observations until the convergence criterion is satisfied or the maximum calibration budget $N_{\max}$ is exhausted.

This ablation separates representation transfer from calibration transfer. The representation is identical in all three deployment strategies; only the information available for estimating the local operational radius changes.

Stage 4 — Multi-unit run-to-failure validation on NASA IMS. A second nominal representation is constructed for the NASA IMS bearing family using simulated healthy observations matched to the dataset sampling and operating characteristics. The representation is trained once and transferred unchanged to each independently monitored bearing.

Each bearing estimates its operational radius from an early-life nominal segment using the same convergence-based calibration procedure adopted in the CWRU experiments. Once convergence is declared, the operational radius is frozen and monitoring continues without further adaptation.

This stage evaluates whether the proposed deployment methodology extends from heterogeneous operating conditions to physically distinct bearing units undergoing naturally evolving degradation.

We report

- (i) observations required for convergence;
- (ii) variability of the converged operational radius across bearings;
- (iii) healthy-period false-alarm rate;
- (iv) first persistent departure after convergence; and
- (v) alarm lead time relative to the documented end of the experiment.

Stage 5 — Edge deployment cost. We report model size, parameter count, CPU inference latency for one 1200-sample window, local calibration memory requirements, and the reduction in transmitted data obtained when only anomaly scores and selected anomalous windows are communicated instead of the continuous raw signal.

5 Results

The results are organized around three complementary deployment questions. First, the CWRU experiments establish a favorable but operationally non-scalable upper reference in which an independent autoencoder, preprocessing transformation, and complete-data calibration are available for each operating condition. Second, they quantify whether one simulator-trained

shared representation, combined only with convergence-based local commissioning, can recover a comparable diagnostic operating point without field-model retraining. Third, the NASA IMS run-to-failure experiment evaluates the complete deployment methodology on independently monitored physical assets undergoing naturally evolving degradation.

The CWRU evaluation considers four operating conditions corresponding to 1730, 1750, 1772, and 1797 RPM. Across these conditions, the experiments contain 1,414 healthy windows, 1,621 ball-fault windows, 1,621 inner-race-fault windows, and 2,846 outer-race-fault windows. All windows contain 1,200 samples and are subjected to the same evaluation protocol, while the representation and calibration information available to each deployment strategy differ according to the experiment.

The NASA IMS evaluation considers four independently monitored bearings from Experiment 2, each represented by 984 chronologically ordered recordings. Unlike the controlled CWRU fault-window comparison, this experiment evaluates whether locally commissioned operational radii remain diagnostically meaningful throughout complete run-to-failure trajectories.

5.1 Ideal asset-specific reference

Table 1 establishes the practical upper reference obtained when independent autoencoders are trained and calibrated for each CWRU operating condition using the complete healthy dataset available for that condition. These experiments constitute an oracle reference rather than a deployable fleet architecture: they assume complete nominal data, an independently fitted preprocessing transformation, and separate representation learning and calibration for every operating environment.

Each asset-specific autoencoder uses the empirical 98th percentile of its own healthy reconstruction-error distribution as its decision boundary. Fault observations are never used for representation learning or threshold selection.

Table 1: Ideal asset-specific autoencoder performance on CWRU. Each autoencoder is trained and calibrated independently using the complete healthy dataset of the corresponding operating condition. FAR denotes healthy false-alarm rate and DR denotes fault detection rate.

RPM	Threshold	Healthy FAR (%)	Ball DR (%)	Inner-race DR (%)	Outer-race DR (%)
1730	0.6485	2.23	100.00	100.00	100.00
1750	0.7590	2.23	100.00	100.00	100.00
1772	0.7248	2.23	100.00	100.00	100.00
1797	0.4591	2.46	100.00	100.00	100.00
Mean	–	2.29 ± 0.12	100.00 ± 0.00	100.00 ± 0.00	100.00 ± 0.00

The asset-specific autoencoders achieve 100% detection for all evaluated fault families under every operating condition while reproducing the expected approximately 2% healthy false-alarm rate induced by empirical p98 calibration. This operating point serves as the reference against which the deployment cost of the proposed methodology is quantified in the following sections.

The absolute reconstruction-error threshold differs substantially across operating conditions, ranging from approximately 0.46 to 0.76. This variation is important for the deployment problem addressed in this paper: even when independently trained models produce identical diagnostic performance, their numerical score scales are not interchangeable.

The asset-specific results therefore define a favorable but operationally non-scalable reference. Each deployment environment receives its own preprocessing transformation, representation model, and complete-data calibration. The subsequent experiments quantify the operational cost of replacing this ideal deployment with one simulator-trained shared representation commissioned independently at each asset through local operational-radius estimation, without any field-model

training.

5.2 Repeated-run stability of the asset-specific autoencoder

To assess the effect of stochastic optimization, the complete asset-specific training and evaluation procedure was repeated five times using independent random initializations. The architecture, preprocessing, datasets, training protocol, and empirical p98 calibration remained unchanged.

Across all operating conditions, healthy false-alarm rate and the detection rates for ball-, inner-race-, and outer-race faults remained unchanged across repetitions. Only the numerical reconstruction threshold varied, as summarized in Table 2.

RPM	Reconstruction threshold
1730	0.649 ± 0.021
1750	0.759 ± 0.227
1772	0.725 ± 0.148
1797	0.459 ± 0.026

Table 2: Stability of the reconstruction threshold

Although the detector was retrained five times on exactly the same healthy dataset, the resulting reconstruction thresholds differed because of stochastic optimization, while the corresponding diagnostic operating points remained unchanged.

This observation has an important implication for fleet deployment. Diagnostic performance is associated with the learned nominal representation rather than with a unique numerical residual scale. Different optimization trajectories produce diagnostically equivalent representations while assigning different absolute reconstruction values to them. Consequently, transferring an absolute threshold between models is generally not meaningful, whereas transferring the shared nominal representation and allowing each deployed asset to identify its own local operational radius is consistent with the observed behaviour.

5.3 Direct simulation-to-field deployment

The shared autoencoder is trained once using simulated healthy observations and transferred unchanged to all four CWRU operating conditions. No CWRU observation is used to retrain, fine-tune, or otherwise modify either the representation or its preprocessing transformation.

Three calibration conditions are evaluated:

1. **Direct threshold transfer**, in which the absolute empirical p98 residual threshold estimated from simulated nominal data is applied unchanged to each field operating condition;
2. **Field oracle**, in which the empirical p98 operational radius is computed retrospectively from the complete healthy field dataset of each operating condition, without retraining the shared representation;
3. **Proposed online calibration**, in which the local empirical p98 operational radius is estimated sequentially from the commissioning residual stream, assumed nominal during initialization. Convergence is evaluated only after a minimum calibration size $N_{\min} = 40$, and calibration continues until the stability criterion is satisfied or the maximum budget $N_{\max} = 200$ is reached.

Once convergence is detected, the corresponding operational radius is frozen and used for subsequent monitoring. The effective calibration duration is therefore determined independently for each operating condition rather than fixed in advance.

Table 3 reports the complete CWRU deployment results.

Table 3: Simulation-to-field deployment of the shared autoencoder on CWRU. FAR denotes healthy false-alarm rate and DR denotes fault detection rate. The field oracle uses every available healthy window, whereas the proposed online procedure stops automatically when the empirical p98 radius satisfies the convergence criterion. Healthy FAR is calculated over the complete healthy dataset to permit direct comparison between calibration strategies. A dash indicates that sequential calibration is not applicable.

RPM	Calibration	τ	Healthy FAR (%)	Ball DR (%)	Inner DR (%)	Outer DR (%)	Healthy windows	Calibration windows
1730	Direct threshold transfer	0.1027	100.00	100.00	100.00	100.00	404	–
1730	Field oracle	0.3578	2.23	99.75	100.00	99.86	404	–
1730	Proposed online calibration	0.3520	3.71	99.75	100.00	100.00	404	49
1750	Direct threshold transfer	0.1027	100.00	100.00	100.00	100.00	404	–
1750	Field oracle	0.3434	2.23	100.00	100.00	100.00	404	–
1750	Proposed online calibration	0.3440	1.98	100.00	100.00	100.00	404	62
1772	Direct threshold transfer	0.1027	100.00	100.00	100.00	100.00	403	–
1772	Field oracle	0.3512	2.23	100.00	100.00	100.00	403	–
1772	Proposed online calibration	0.3550	1.24	100.00	100.00	100.00	403	58
1797	Direct threshold transfer	0.1027	100.00	100.00	100.00	100.00	203	–
1797	Field oracle	0.3886	2.46	98.03	100.00	100.00	203	–
1797	Proposed online calibration	0.3869	3.45	98.52	100.00	100.00	203	49

Direct transfer of the absolute simulator threshold fails under every operating condition. Although it produces an apparent 100% detection rate for all evaluated fault families, it simultaneously classifies every healthy field window as anomalous. The absolute residual scale learned in simulation therefore does not transfer directly to the field.

In contrast, local calibration restores an operationally meaningful decision boundary without modifying the shared representation. The complete-data field oracle reproduces the expected approximately 2% healthy false-alarm rate and achieves 99.45% average ball-fault detection, 100% inner-race detection, and 99.96% outer-race detection.

The proposed online calibration reaches a closely comparable operating point using only the observations available during commissioning. All four operating conditions satisfy the convergence criterion before the maximum calibration budget is exhausted, requiring 49, 62, 58, and 49 windows for 1730, 1750, 1772, and 1797 RPM, respectively. The mean calibration duration is therefore 54.5 ± 6.6 windows.

Across the four conditions, the proposed procedure obtains a mean healthy FAR of $2.60 \pm 1.18\%$, 99.57% average ball-fault detection, 100% inner-race detection, and 100% outer-race detection. These results are close to the complete-data oracle despite requiring neither complete prior knowledge of the healthy field distribution nor independent representation learning for each operating condition.

The central result is therefore not that the online procedure outperforms the complete-data oracle. Rather, it approximates the oracle operating point through a short, automatically terminated commissioning phase while retaining one fixed fleet-wide representation.

5.4 Cost of scalable deployment

Table 4 compares the proposed shared deployment with an ideal but operationally non-scalable asset-specific reference. The comparison intentionally favors the baseline: the oracle detector is trained and calibrated using the complete healthy dataset of each operating condition, whereas the proposed methodology performs no field-model training and estimates the local operational radius through convergence-based commissioning, requiring only 49–62 healthy observations before deployment.

Table 4: Ideal but non-scalable asset-specific deployment versus the proposed fleet-scalable deployment. Values are mean $\pm$ standard deviation across the four CWRU operating conditions.

Deployment strategy	Healthy FAR (%)	Ball DR (%)	Inner DR (%)	Outer DR (%)	Field model training	Healthy field data used
Asset-specific AE oracle	2.29 ± 0.12	100.00 ± 0.00	100.00 ± 0.00	100.00 ± 0.00	One model per condition	Complete dataset
Shared AE with field-oracle radius	2.29 ± 0.12	99.20 ± 1.60	100.00 ± 0.00	100.00 ± 0.00	None	Complete dataset
Proposed shared AE with online radius	3.09 ± 2.22	99.69 ± 0.62	100.00 ± 0.00	100.00 ± 0.00	None	100-window warm-up

Replacing four independently trained autoencoders with one simulator-trained shared model produces no observed loss for inner- or outer-race detection and less than one percentage point of average difference for ball-fault detection. The principal cost appears in healthy false alarms, whose mean increases from 2.29% to 3.09%. Most of this increase is concentrated in the 1797 RPM condition (5.91%), whereas the remaining conditions range between 0.50% and 2.97%.

These results support the central deployment claim of this paper. A single simulator-trained canonical representation, combined with short local operational-radius calibration, achieves near-oracle detection performance without asset-specific representation learning, substantially reducing deployment complexity.

Table 5: Agreement between the converged online operational radius and the complete-data field-oracle p98 reference. Relative difference is computed as $100|\tau_{\text{online}} - \tau_{\text{oracle}}|/\tau_{\text{oracle}}$. Post-calibration FAR is evaluated only on healthy windows occurring after the declared convergence point.

RPM	τ_{oracle}	τ_{online}	Relative difference (%)	Calibration windows	Post-calibration FAR (%)
1730	0.3578	0.3520	1.64	49	3.94
1750	0.3434	0.3440	0.19	62	1.75
1772	0.3512	0.3550	1.06	58	0.87
1797	0.3886	0.3869	0.44	49	3.90
Mean	—	—	0.83	54.5 ± 6.6	2.62

5.5 Operational-radius convergence

Figure 2 illustrates the sequential estimation of the local operational radius under the fixed shared representation. During commissioning, the simulator contributes neither an absolute reconstruction-error threshold nor a standardized residual scale. The only fleet-wide calibration quantity is the target nominal coverage, $1 - \alpha = 0.98$.

For operating condition k , the operational radius available after observing t commissioning residuals is estimated directly as

$$\tau_k(t) = \hat{Q}_{k,t}(0.98), \quad (3)$$

where $\hat{Q}_{k,t}(0.98)$ denotes the empirical 98th percentile of the local reconstruction residuals observed up to commissioning instant t . The shared representation remains unchanged throughout this process; only the asset-specific decision boundary is estimated locally.

Convergence is evaluated after a minimum calibration period $N_{\text{min}} = 40$. Using the criterion defined in Eq. (??), calibration terminates when the relative change in $\tau_k(t)$ remains below the prescribed tolerance for the required number of consecutive updates. If convergence is not observed, calibration continues until the maximum commissioning budget $N_{\text{max}} = 200$ is reached. An asset that does not converge within this interval is not released into active monitoring and is instead assigned a calibration-failure status.

For retrospective comparison, the complete-data field-oracle radius is

$$\tau_k^{\text{oracle}} = \hat{Q}_{k,\text{all}}(0.98), \quad (4)$$

where the quantile is computed using every available healthy field window for operating condition k . This quantity is unavailable during commissioning and is used only to assess how closely the online procedure approximates the operating point obtained with complete nominal knowledge.

All four CWRU operating conditions satisfy the convergence criterion before the maximum calibration budget is exhausted. Convergence is detected after 49, 62, 58, and 49 windows for 1730, 1750, 1772, and 1797 RPM, respectively, corresponding to a mean commissioning duration of 54.5 ± 6.6 windows. At convergence, the resulting operational radii differ from their complete-data field-oracle references by only 1.64%, 0.19%, 1.06%, and 0.44%, respectively.




Figure 2: Sequential estimation of the empirical p98 operational radius during CWRU commissioning. For each operating condition, the blue segment shows the online estimate $\tau_k(t) = \hat{Q}_{k,t}(0.98)$ until the convergence criterion is satisfied. The red dotted line marks the declared convergence point, at which the operational radius is frozen and active monitoring begins. The gray segment is a retrospective continuation obtained by incorporating additional observations known from the benchmark to be healthy; it is shown only to characterize the subsequent behaviour of the estimator and is not used to modify the deployed decision boundary. The horizontal dashed line indicates the complete-data field-oracle p98 reference.

The retrospective trajectories show that an empirical upper quantile can continue to change after an apparently stable interval when additional high-residual nominal observations arrive. This behaviour is expected because, during short commissioning periods, the empirical p98 estimate is determined by only a small number of observations in the upper tail. The post-convergence trajectories therefore do not represent continued adaptation of the deployed detector; they show how the estimator would have evolved had calibration remained active.

The operationally relevant result is the close agreement between the frozen online radius and the complete-data oracle. Despite using only 49–62 commissioning observations, the online procedure recovers an acceptance boundary within 2% of the radius obtained retrospectively from all available healthy observations. The resulting post-calibration healthy false-alarm rates range from 0.87% to 3.94%, remaining close to the nominal 2% operating point despite the substantially reduced calibration information.

The purpose of the convergence criterion is not to prove that the population 98th per-

centile has been recovered exactly. It provides an operational stopping rule indicating that the local radius has remained sufficiently stable for deployment under the selected parameters ($N_{\min}$, $N_{\max}$, L , ε_C , m). These parameters may be adjusted to reflect the commissioning duration and false-alarm tolerance of a particular application.

The present experiment assumes that commissioning observations are representative of nominal operation. Failure to obtain a stable operational radius within $N_{\max}$ may indicate insufficient commissioning data, non-stationary operating conditions, or an asset whose initial behaviour is inconsistent with the assumed nominal state. Distinguishing among these causes is outside the present evaluation, but the calibration-failure status provides a conservative mechanism that prevents an unconverged asset from entering active monitoring.

5.6 Robustness to local initialization

A practical deployment question is whether convergence-based local calibration depends critically on the particular nominal observations encountered during commissioning. To evaluate this sensitivity, the shared simulator-trained representation, preprocessing transformation, target coverage, and convergence parameters were kept fixed throughout the experiment. Only the nominal observations available during commissioning, together with their arrival order, were varied.

For each CWRU operating condition, five independent commissioning sequences were generated by randomly permuting the available healthy observations. Calibration was allowed to start after $N_{\min} = 40$ observations and continued until the empirical p98 operational radius satisfied the fixed convergence criterion or the maximum commissioning budget $N_{\max} = 200$ was exhausted. Once convergence was declared, the estimated operational radius was frozen and subsequently used for fault detection. Healthy false-alarm rate was evaluated only on healthy observations excluded from calibration, while all fault classes were evaluated using the same shared representation without retraining.

Table 6 summarizes the results obtained across the five independent commissioning sequences.

Table 6: Robustness of convergence-based local calibration to commissioning sample variability. The shared simulator-trained representation and all calibration parameters remained fixed; only the healthy commissioning observations and their arrival order were varied. Values are reported as mean $\pm$ standard deviation over five independent commissioning realizations. FAR is evaluated only on healthy observations excluded from calibration.

RPM	τ_{online}	Relative oracle error (%)	Post-calibration FAR (%)	Ball DR (%)	Inner DR (%)	Outer DR (%)	Calibration windows
1730	0.3486 $\pm$ 0.0023	2.59 $\pm$ 0.63	5.38 $\pm$ 1.39	99.95 $\pm$ 0.11	100.00 $\pm$ 0.00	100.00 $\pm$ 0.00	55.6 $\pm$ 8.2
1750	0.3451 $\pm$ 0.0111	2.53 $\pm$ 1.65	4.23 $\pm$ 5.63	99.90 $\pm$ 0.14	100.00 $\pm$ 0.00	100.00 $\pm$ 0.00	57.4 $\pm$ 7.9
1772	0.3511 $\pm$ 0.0086	1.82 $\pm$ 1.37	2.54 $\pm$ 2.18	100.00 $\pm$ 0.00	100.00 $\pm$ 0.00	99.97 $\pm$ 0.06	54.4 $\pm$ 5.9
1797	0.3915 $\pm$ 0.0065	1.40 $\pm$ 1.02	2.21 $\pm$ 2.23	97.98 $\pm$ 0.71	100.00 $\pm$ 0.00	99.97 $\pm$ 0.06	64.4 $\pm$ 26.0

All 20 commissioning runs satisfied the convergence criterion before reaching the maximum calibration budget, yielding a calibration success rate of 100%. Across all operating conditions, the estimated operational radius remained within $2.08 \pm 0.71\%$ of the complete-data oracle, while the fleet-averaged post-calibration healthy false-alarm rate was $3.59 \pm 1.59\%$. Fault-detection performance remained essentially unchanged throughout all commissioning realizations, indicating that the shared representation is insensitive to realistic variability in the local initialization process.

These experiments evaluate the robustness of the complete commissioning procedure rather than the variability of a fixed-size percentile estimate. The representation, preprocessing pipeline, target coverage, convergence criterion, and all deployment parameters remained unchanged. Only the local healthy observations available during commissioning and their arrival order were modified.

Some variability is expected because the empirical p98 operational radius is estimated from a relatively small number of upper-tail observations. Consequently, different commissioning

samples naturally produce slightly different estimates of the local operational radius. These differences translate into modest variations in the healthy false-alarm rate while leaving the resulting diagnostic operating point essentially unchanged.

The largest healthy false-alarm rate observed during the repeated commissioning experiments was 13.8%. However, this isolated realization did not alter the deployment outcome: convergence was still achieved, the estimated operational radius remained within 4% of the oracle estimate, and fault-detection performance remained essentially unchanged.

The observed variability therefore reflects finite-sample uncertainty in estimating an extreme empirical quantile rather than instability of the proposed deployment methodology. From the quotient-space formulation, this behaviour is expected: the operational radius estimates the local neighbourhood of the nominal manifold rather than representing an arbitrary decision threshold. Consequently, different healthy commissioning samples should produce slightly different neighbourhood estimates while preserving the underlying nominal representation.

Together with the oracle-agreement results of Table 5, these experiments establish that convergence-based commissioning reliably recovers the operational neighbourhood of the shared nominal manifold under realistic initialization variability. The oracle-agreement experiment evaluates chronological calibration against the complete-data reference, whereas the present experiment demonstrates that essentially the same neighbourhood is recovered under different commissioning realizations without modifying the shared representation.

The independent NASA IMS validation presented in Section 5.7 extends this conclusion beyond controlled operating conditions. There, independently commissioned physical assets converged to distinct local operational radii while subsequently exhibiting persistent departures from their respective nominal neighbourhoods during naturally evolving degradation. Taken together, the CWRU and NASA experiments support the central deployment hypothesis of this work: the transferable object is the shared nominal representation, whereas each asset identifies only its own local operational radius during commissioning.

5.7 Multi-unit run-to-failure deployment on NASA IMS

The CWRU experiments demonstrate simulation-to-field transfer under controlled operating conditions. The NASA IMS experiments evaluate the complete deployment methodology under realistic fleet conditions, where independently operating physical assets undergo naturally evolving degradation over their remaining service life.

Each NASA bearing is treated as an independent monitored asset. The simulator-trained autoencoder and preprocessing pipeline are transferred unchanged, and no NASA observations are used for representation learning, retraining, or parameter adaptation. The only asset-specific operation is the estimation of the local operational radius from only the early-life recordings selected for commissioning and their arrival order were varied.

Table 7 summarizes the deployment results. Operational-radius calibration converged successfully for all four bearings after only 49–57 healthy recordings. Although the estimated operational radii differ substantially across bearings, every asset uses exactly the same shared representation, the same target nominal coverage (p98), and the same convergence criterion without bearing-specific retraining.

Following commissioning, all four bearings subsequently departed persistently from their estimated nominal neighbourhood. The learned operational radius remains predictive throughout the subsequent degradation trajectory, with a mean post-departure exceedance rate of 94.6% and lead times ranging from 14.5 h to 74.0 h before the documented end of the experiment.

Table 7: NASA IMS deployment results using one shared simulator-trained representation and convergence-based local operational-radius estimation.

Bearing	Calibration (windows)	τ_k	Relative error (%)	Post-departure exceedance (%)	Lead time (h)
1	49	0.010680	0.10	99.12	74.00
2	57	0.014950	0.35	86.42	38.83
3	49	0.018930	2.87	95.88	14.50
4	49	0.006672	0.18	97.02	54.33
Mean	51	–	0.87	94.61	45.42

These results demonstrate that the transferable engineering artefact is the shared nominal representation rather than an asset-specific detector. The same representation identifies substantially different operational radii for different physical assets while requiring no representation retraining and no threshold engineering. Once identified during commissioning, each operational radius remains predictive throughout the subsequent degradation trajectory. The same learned representation produces substantially different operational radii for different physical assets while requiring neither representation retraining nor threshold engineering. Once estimated during commissioning, the local operational radius continues to separate nominal from degrading behaviour throughout the remaining lifetime of the asset.

5.7.1 Stability of local operational-radius estimation

A practical deployment methodology must be robust to the particular healthy observations collected during commissioning. To evaluate this property, the shared representation was kept completely fixed while only the chronological commissioning sample used to estimate the local operational radius was varied.

For each NASA bearing, five independent commissioning realizations were generated by sampling and permuting the first 200 early-life recordings. No retraining was performed; only the observed nominal residuals differed between runs.

Across twenty independent commissioning experiments (four bearings and five random initialization samples), calibration converged successfully in every case. Table 8 summarizes the resulting variability.

Table 8: Stability of local operational-radius estimation under repeated NASA commissioning experiments using one fixed shared representation.

Bearing	τ_{mean}	τ_{std}	Relative error (%)	Healthy FAR (%)	Calibration (windows)	Success (%)
1	0.010590	0.000074	0.73 ± 0.70	5.84 ± 3.79	49.2 ± 0.4	100
2	0.014873	0.000084	0.36 ± 0.44	3.71 ± 3.33	49.0 ± 0.0	100
3	0.018151	0.000383	1.58 ± 1.88	4.93 ± 4.22	52.2 ± 7.2	100
4	0.006640	0.000033	0.66 ± 0.49	5.45 ± 2.39	49.6 ± 1.3	100
Mean	–	–	0.83	4.98	50.0	100

These experiments demonstrate that the operational radius is simultaneously asset-specific and statistically stable. Different physical units identify different local neighbourhoods of the same shared nominal representation, whereas repeated commissioning of the same unit produces only minor variation in the estimated radius. The transferable engineering artefact is therefore the shared nominal representation itself, while commissioning consists only of identifying the local operational neighbourhood associated with each physical asset.

5.8 Edge deployment implications

The shared model contains 847,601 trainable parameters and is trained entirely offline. Deployment at the monitored asset requires only forward inference, a small residual buffer, estimation

of an empirical quantile, and the current convergence state. No gradient computation, model optimization, fault labels, or asset-specific representation learning are required after installation.

On the evaluation platform, inference over a 1,200-sample window requires a median CPU latency of only 6.26 ms (95th percentile: 6.62 ms), substantially below the 100 ms acquisition interval corresponding to the evaluated windows. Consequently, the proposed deployment operates comfortably in real time while leaving ample computational margin for additional edge tasks such as communication, diagnostics, or signal preprocessing.

The local deployment state is also extremely compact. Besides the shared model itself, each asset stores only a buffer of the most recent 200 reconstruction residuals during commissioning, requiring less than 1 kB of asset-specific state. The only fleet-wide artifact is the shared nominal representation, which is trained once and distributed to every operational unit.

Table 9: Edge-deployment footprint of the proposed shared nominal representation.

Metric	Value	Metric	Value
Model parameters	847,601	Residual-buffer memory	800 bytes
Raw float32 size	3.39 MB	Total calibration state < 1 kB	
Serialized (.keras)	10.24 MB	Per-asset model training	None
CPU latency (median)	6.26 ms	Fault labels required	None
CPU latency (p95)	6.62 ms	Local residual buffer	200 samples

6 Mathematical Scope and Deployment Guarantees

This section states, without full proof, the minimal functional-analytic scaffolding justifying why one shared canonical representation and independently estimated local operational radii are mathematically compatible; complete derivations are given in our companion technical reports Di Santi 2026a,b.

Scope. Jointly recovering a nuisance-invariant representation *and* calibrating per-asset decision boundaries across a heterogeneous fleet is combinatorially hard. This paper decouples that problem into (i) a *global* inverse problem, solved once offline, recovering a canonical representation whose nominal orbit is relatively compact (Proposition 1); and (ii) a *local* inverse problem, solved independently by each asset as an $O(1)$ -per-observation empirical-quantile estimate.

6.1 Observation model

Let $X \subset \mathbb{R}^n$ denote observation windows, $\theta \in \Theta$ the latent diagnostic state, and $\eta \in E$ nuisance variables (sensor coupling, mounting, operating point, noise), with $x = F(\theta, \eta)$. The group G_{nuis} acts on E while preserving θ ; observations are diagnostically equivalent, $x \sim x'$, iff $x' = F(\theta, g \cdot \eta)$ for some $g \in G_{\text{nuis}}$.

6.2 Compactness and factorization

Proposition 1 (Relative compactness of the nominal orbit). *If E is compact and $\eta \mapsto F(\theta_N, \eta)$ is uniformly bounded and equicontinuous on E , then $O_N = \{F(\theta_N, \eta) : \eta \in E\}$ is relatively compact in $(C(X), \|\cdot\|_\infty)$.*

Proof sketch. By Arzelà–Ascoli Rudin 1976, boundedness and equicontinuity on compact E give relative compactness of O_N ; the equicontinuity modulus for the Stage 1 simulator is verified in Di Santi 2026b. $\square$

Because the nominal orbit is compact, τ_k^* only accounts for asset-specific scale around a fixed structure rather than searching an unbounded space, explaining the rapid convergence observed empirically (51.0 ± 12.6 windows on CWRU, Section 5.5).

Proposition 2 (Factorization through the nuisance quotient). *Let $\pi : X \rightarrow X/G_{\text{nuis}}$ be the quotient projection. If $d : X \rightarrow \mathbb{R}$ is G_{nuis} -invariant, then d is constant on fibers and factors uniquely as $d = \bar{d} \circ \pi$.*

Proof sketch. Invariance makes d constant on fibers, so $\bar{d}([x]) := d(x)$ is well defined by the universal property of the quotient map Di Santi 2026b. $\square$

Together, these results justify treating diagnosis as membership testing around a fixed canonical representation: $\bar{d}$ (equivalently O_N) is recovered once, globally, while each asset k solves only a well-posed local calibration, estimating τ_k^* for coverage $1 - \alpha$ (Eq. (??)) from its own residuals. Empirically, this is reflected in the same fixed representation transferring across four CWRU conditions with only $\tau_k(t)$ re-identified (99.69% average ball-fault, 100% inner/outer-race detection), and extending to physically distinct NASA IMS bearings with 83.25 h average lead time.

7 Discussion

Calibration should not be interpreted as parameter tuning. Rather, it identifies an unknown physical property of the monitored asset: its nominal operational radius.

The principal result of this work is not that a Conv1D autoencoder can separate CWRU fault classes, a capability already demonstrated extensively in the literature. The contribution lies instead in the deployment methodology: a single canonical nominal representation can be shared across heterogeneous assets, while each physical unit independently estimates its own operational radius without retraining the representation.

From a procurement perspective, this separation gives fleet operators a concrete basis for specifying deployment requirements: rather than accepting per-asset models developed ad hoc by each vendor, an operator can require a single validated, versioned nominal representation together with documented local calibration behaviour, independent of which organization, component OEM, third-party analytics provider, or in-house engineering team, produces and maintains it.

The repeated-run and initialization-robustness experiments jointly support this claim from two complementary directions. Independent random initializations converge to slightly different reconstruction-error scales, yet healthy false-alarm rate and fault detection remain essentially unchanged after empirical calibration. Conversely, keeping the learned representation fixed while varying only the commissioning sample produces only modest variation in the estimated operational radius. The transferable object is therefore not an absolute reconstruction-error threshold but a shared canonical representation together with a target nominal coverage. The calibration-transfer experiments demonstrate that the representation transfers successfully from simulation to field, whereas the absolute residual scale does not.

Identifying the shared representation as the transferable object also changes the engineering organization of fleet deployment. Rather than developing and maintaining an independent representation model for each physical asset, these activities can be centralized, for example by the original equipment manufacturer (OEM), while each deployed asset performs only a lightweight local calibration during commissioning. Computational scalability is correspondingly decoupled from fleet size: growth increases the number of monitored assets, not the number of engineered models, since adding a unit requires only distributing the shared representation and estimating its local operational radius.

Beyond commissioning, the NASA IMS experiments indicate that the estimated operational radius remains predictive throughout the subsequent degradation trajectory. Across all evaluated

bearings, calibration converged after approximately fifty healthy observations, while persistent degradation subsequently evolved outside the learned nominal neighbourhood for the vast majority of the remaining asset lifetime. Although the four bearings estimated substantially different operational radii, all of them used the same shared representation, the same target nominal coverage (p98), and the same convergence criterion, without bearing-specific retraining.

This behaviour is consistent with the broader theoretical framework developed in our previous work, which formulates industrial diagnosis as an inverse problem over quotient spaces Di Santi 2025, 2026a,b. Within that formulation, the operational radius is not merely a tuned threshold but an estimate of the local neighbourhood of the nominal manifold. Consequently, once calibration has converged, trajectories that remain nominal are expected to remain inside this neighbourhood, whereas persistent degradation should naturally evolve outside it. The NASA IMS experiments provide empirical evidence consistent with this theoretical prediction.

This perspective shifts commissioning from threshold tuning to manifold identification. Rather than adjusting detector parameters for each individual asset, deployment consists of estimating the local operational neighbourhood associated with a shared nominal representation. The learned representation becomes the transferable engineering artefact, while commissioning reduces to identifying the operational radius of each physical asset.

Operational-radius convergence also clarifies the role of adaptation: its purpose is to identify the asset during commissioning, not to track its aging indefinitely. Once τ_k^* has been established, continued movement of the anomaly score constitutes diagnostic evidence rather than input for further adaptation. Initialization can therefore proceed until a predefined stability criterion is satisfied, after which the operational radius is frozen or only tightly constrained. Together with the edge footprint reported in Section 5.8, this deployment model enables day-one monitoring without fault labels, isolates asset-specific adaptation to local calibration, and requires only lightweight edge computation.

Conventional CBM deployments typically require asset-specific engineering of both models and alarm thresholds. The present results suggest that this engineering effort can instead be replaced by a standardized commissioning procedure in which every asset estimates its own operational radius while reusing a common nominal representation.

Remark 1. *The variability observed for the 1750 RPM commissioning samples is not specific to the shared-representation architecture. It reflects the finite-sample uncertainty inherent in estimating an upper empirical quantile before the complete nominal distribution is available. An asset-specific detector calibrated from the same limited commissioning sample would be subject to the same statistical limitation; the complete-data oracle avoids it only by assuming retrospective access to all healthy observations.*

The convergence rule adopted in the present experiments is operational rather than universal. Calibration is declared after the relative variation of the empirical p98 radius remains below $\varepsilon_C = 0.01$ over a look-back interval $L = 10$ for $m = 10$ consecutive updates. These parameters define a configurable trade-off between commissioning duration and robustness against transient plateaus in the upper-tail estimate. Requiring more consecutive stable updates, a longer look-back interval, or a tighter tolerance would generally delay deployment while reducing sensitivity to short-term fluctuations; conversely, less restrictive settings would shorten commissioning at the cost of increased threshold variability. The present study does not claim that $(L, \varepsilon_C, m) = (10, 0.01, 10)$ is optimal. Rather, it demonstrates that this single fixed configuration achieved successful convergence for all evaluated CWRU and NASA IMS commissioning experiments while preserving near-oracle diagnostic behaviour. Systematic selection of these parameters according to application-specific false-alarm and commissioning-time requirements remains an important direction for future work.

7.1 Limitations

Several limitations qualify the scope of the present results and should guide their interpretation.

Nominal initialization assumption. The procedure assumes the cold-start observations are representative of healthy operation. If an asset were already degraded at installation, τ_k^* could absorb part of that condition into the nominal baseline; a consistency check against the canonical nominal manifold is a natural extension and is left as future work.

Short-warm-up sensitivity. Because the operational radius is an empirical 98th-percentile estimated from only 100 windows, it can be sensitive to a small number of extreme values, as observed in the 1797 RPM condition (5.91% FAR, roughly twice the other conditions). Longer warm-ups or robust quantile estimators would likely reduce this variability at the cost of a longer cold start.

Stationarity of nominal operation. The convergence criterion assumes an approximately stationary healthy residual distribution during cold start. The retrospective trajectories in Figure 2 show modest post-calibration drift; more variable startup conditions were not evaluated and may benefit from a longer or adaptively triggered calibration window.

Validation domain. The validation uses bearing datasets (CWRU, NASA IMS) rather than instrumented transportation assets. While the deployment principle is architecture- and modality-agnostic by construction (Section 3), its extension to other components remains to be confirmed empirically.

7.2 Future Work

A natural extension relaxes the assumption that assets begin in nominal condition. Because the methodology separates the shared representation from the local radius, future work will investigate validating the initialization phase itself against the canonical nominal manifold, to distinguish genuinely healthy startup from incipient degradation before it is absorbed into the baseline. A second direction concerns progressive degradation: whether the joint evolution of reconstruction residuals and operational-radius dynamics carries predictive information, connecting this methodology with physics-informed health indicators and remaining-useful-life estimation.

Finally, the mathematical interpretation in Di Santi 2026a suggests formalizing the operational radius as an asset-specific neighborhood around the transferable canonical representation, characterizing its invariance properties and the conditions under which local calibration preserves the underlying quotient-space structure.

8 Conclusion

We presented a day-one deployment methodology for nominal condition monitoring across heterogeneous asset fleets, validated on public rolling-element bearing datasets. A single nominal representation is trained once on simulated healthy observations and transferred unchanged to field deployment. Each operational unit then performs only a short unlabeled warm-up to estimate its own operational radius, without modifying or retraining the shared representation.

The experimental results show that a simulator-trained nominal representation transfers successfully to field deployment, whereas an absolute decision threshold does not. A shared target nominal coverage, realized locally through the empirical residual quantile estimated during initialization, provides an effective bridge between the shared representation and field deployment. This deployment strategy preserves detection performance close to that of asset-specific models trained under ideal asset-specific calibration conditions while eliminating per-asset representation learning.

The transferable object is therefore not an absolute anomaly threshold, nor an asset-specific representation, but a shared canonical nominal representation together with a lightweight local

calibration mechanism. The central outcome of this work is consequently not a new anomaly detector, but a scalable deployment principle: learn the canonical nominal representation once, distribute it across the fleet, and let each physical asset identify only its own operational radius during installation.

This separation between representation learning and deployment calibration provides a practical foundation for large-scale edge condition monitoring, where an OEM can centrally produce and maintain the shared representation while each deployed asset performs its own lightweight local calibration automatically, without operator intervention or manual threshold tuning.

Conventional CBM deployments typically require asset-specific engineering of models and alarm thresholds. The present results suggest that this engineering effort can instead be replaced by a standardized commissioning procedure in which every asset estimates its own operational radius while reusing a common nominal representation. The representation is shared. The operational radius is individual.

References

- Bishop, Christopher M. (1994). “Novelty Detection and Neural Network Validation”. In: *IEE Proceedings – Vision, Image and Signal Processing* 141.4, pp. 217–222. DOI: 10.1049/ip-vis:19941330.
- Bronstein, Michael M., Joan Bruna, Taco Cohen, and Petar Veličković (2021). “Geometric Deep Learning: Grids, Groups, Graphs, Geodesics, and Gauges”. In: *arXiv preprint*. arXiv: 2104.13478.
- Carvalho, João B. S., Mengtao Zhang, Robin Geyer, Carlos Cotrini, and Joachim M. Buhmann (2023). “Invariant Anomaly Detection under Distribution Shifts: A Causal Perspective”. In: *Advances in Neural Information Processing Systems*. Vol. 36.
- Chandola, Varun, Arindam Banerjee, and Vipin Kumar (2009). “Anomaly Detection: A Survey”. In: *ACM Computing Surveys* 41.3, 15:1–15:58. DOI: 10.1145/1541880.1541882.
- Chen, [Nombre] et al. (2026). “A digital twin-driven cross-domain adaptation method for bearing intelligent fault diagnosis”. In: *Systems Science & Control Engineering*.
- Di Santi, Eduardo (2025). “The Blueprints of Intelligence: A Functional-Topological Foundation for Perception and Representation”. In: arXiv: 2512.05089.
- (2026a). *Anomaly Detection as an Inverse Problem on Quotient Spaces: Equivalence Relations and Functional Manifolds as a General Diagnostic Framework*. Technical Report. University of Colorado Boulder. URL: <https://scholar.colorado.edu/concern/reports/7h149s02s> (visited on 07/17/2026).
- (2026b). *Self-Supervision, Manifold Saturation, and Emergent Operational Classes: Learning from Scarce Observations as Inverse Recovery of Compact Functional Structure*. Technical Report. University of Colorado Boulder. URL: <https://scholar.colorado.edu/concern/reports/fn107102t> (visited on 07/17/2026).
- Entezami, Mani, Clive Roberts, Paul Weston, Edward Stewart, Arash Amini, and Mayorkinos Papaelias (2020). “Perspectives on railway axle bearing condition monitoring”. In: *Proceedings of the Institution of Mechanical Engineers, Part F: Journal of Rail and Rapid Transit*.
- Hinton, Geoffrey E. and Ruslan R. Salakhutdinov (2006). “Reducing the Dimensionality of Data with Neural Networks”. In: *Science* 313.5786, pp. 504–507. DOI: 10.1126/science.1127647.
- Hue, Nguyen Thi, Nguyen Duc Thuan, and Hoang Si Hong (2025). “Knowledge-driven subdomain adaptation for bearing fault diagnosis”. In: *Journal of Vibration and Control*.
- Jardine, Andrew K. S., Daming Lin, and Dragan Banjevic (2006). “A Review on Machinery Diagnostics and Prognostics Implementing Condition-Based Maintenance”. In: *Mechanical Systems and Signal Processing* 20.7, pp. 1483–1510. DOI: 10.1016/j.ymssp.2005.09.012.

- Randall, Robert B. and Jérôme Antoni (2011). “Rolling Element Bearing Diagnostics — A Tutorial”. In: *Mechanical Systems and Signal Processing* 25.2, pp. 485–520. DOI: 10.1016/j.ymssp.2010.07.017.
- Rudin, Walter (1976). *Principles of Mathematical Analysis*. 3rd. McGraw-Hill.
- Sakurada, Mayu and Takehisa Yairi (2014). “Anomaly Detection Using Autoencoders with Nonlinear Dimensionality Reduction”. In: *Proceedings of the 2nd Workshop on Machine Learning for Sensory Data Analysis*. ACM, pp. 4–11. DOI: 10.1145/2689746.2689747.
- Schölkopf, Bernhard, Robert C. Williamson, Alex J. Smola, John Shawe-Taylor, and John C. Platt (1999). “Support Vector Method for Novelty Detection”. In: *Advances in Neural Information Processing Systems*. Vol. 12.
- Smith, Wade A. and Robert B. Randall (2015). “Rolling Element Bearing Diagnostics Using the Case Western Reserve University Data: A Benchmark Study”. In: *Mechanical Systems and Signal Processing* 64–65, pp. 100–131. DOI: 10.1016/j.ymssp.2015.04.021.
- Wang, B. et al. (2026). “Fleet-based transfer learning for anomaly detection in [machinery]”. In: *Reliability Engineering & System Safety*.
- Zhao, Rui, Ruqiang Yan, Zhenghua Chen, Kezhi Mao, Peng Wang, and Robert X. Gao (2019). “Deep Learning and Its Applications to Machine Health Monitoring”. In: *Mechanical Systems and Signal Processing* 115, pp. 213–237. DOI: 10.1016/j.ymssp.2018.05.050.